

**CINEAI LAB**

# **Experimental Foundations for Open AI-Native Cinema**

*A consolidated manuscript on creative experiments, sustained human-AI  
collaboration, open infrastructure, global technology watch and governed  
innovation*

**Renato Barrios · CineAI Lab**

English international working edition · Version 0.9 · August 2026

*This edition consolidates the first working draft into one coherent text. It does not reproduce protected screenplays,  
dialogue, endings or unpublished character arcs. CineAI Lab has been developed through sustained iterative  
collaboration with ChatGPT, while protected narrative authorship and final creative decisions remain human.*

## Editorial note

This manuscript replaces the earlier sequence of overlapping volumes with a single, concise account of the CineAI Lab experiment. Repeated lesson tables, duplicated conclusions and detailed scenario summaries have been removed. The document now focuses on the effects sought in experimentation, the evidence obtained, the limitations observed and the case for an open production environment.

The original narrative works remain protected. Their titles may appear publicly, but this manuscript does not disclose complete plots, dialogue, endings or private character histories. The demonstrators are described only through the cinematic capabilities they require: sustained human performance, multilingual speech, non-human movement, nested memory, ensemble continuity, symbolic objects and selective revision.

The experiments are not presented as finished films. Their purpose is to show what is already possible, what still fails, and which forms of integration could turn isolated generative effects into a controllable audiovisual process.

## Why this manuscript is published in English

This manuscript is published in English as its sole canonical edition. The choice is pragmatic rather than cultural: English currently functions as the principal shared working language across artificial-intelligence research, open-source software, model documentation, international audiovisual technology, technical conferences and cross-border development communities. Publishing in that language gives the work the broadest direct access to researchers, developers, producers, rights specialists and potential collaborators who do not share a national language.

Maintaining one authoritative text also protects continuity. Parallel translations would require repeated legal, technical and editorial revision, could introduce differences between versions, and would consume resources better used for experimentation. The multilingual CineAI Lab website provides orientation in English, Spanish, Chinese and French, while the downloadable manuscript remains a single English reference that can be cited, revised and versioned without translation drift.

This decision does not diminish the French origin of the project or the importance of francophone cultural debate. France and Europe remain central subjects of the manuscript, particularly in the sections on public investment, Mistral AI, cultural institutions and adaptive governance. English is used here as an international bridge, not as a replacement for the languages in which the works themselves may be conceived or produced.

## Executive summary

CineAI Lab began from a visible contradiction. Artificial intelligence could already create striking still images, persuasive fragments of voice and short video clips, yet the transition from those fragments to coherent cinema remained unstable. Characters changed across shots. Camera movement was mistaken for acting. Speech could be synchronized without becoming a performance. Animal movement lost weight and intention. A revised shot could damage everything around it.

The first cycle therefore became an experiment in the conditions of filmmaking rather than a race to produce the longest possible AI-generated sequence. Original narrative projects served as demanding test beds, but the protected stories were never the product of the benchmark. The films judged the technology, not the reverse.

The work established several practical findings:

- The laboratory itself became a human–AI workflow: ChatGPT helped connect narrative development, experimental design, technical troubleshooting, documentation and technology watch across otherwise fragmented disciplines.
- Still-image quality is a major preproduction advantage, but it does not predict temporal coherence.
- A valid MP4, a moving mouth or a camera push is a technical output, not proof of cinematic performance.

- Voice, lips, gaze, breathing, gesture, listening and timing form one coupled performance system.
- Every transfer between tools can alter identity, colour, duration, resolution, rights metadata or emotional intent.
- Short, controlled shots reveal more than ambitious end-to-end demonstrations whose failures cannot be diagnosed.
- Commercial suites reduce access friction, but they do not automatically provide narrative memory, selective editing or production independence.
- Open-source repositories are essential, but a catalogue of components is not yet an accessible film environment.

The next objective is an open, modular and free core environment that preserves project memory, identity, performance direction, version history and provenance while allowing different generation engines to be exchanged. Free does not mean that computation has no cost. It means that the orchestration layer, project format, local tools and essential workflows should not require permanent subscription or vendor lock-in.

## Contents

1. Purpose of the experiment
    - 1.1 Human–AI collaboration: ChatGPT as a working partner
  2. Creative demonstrators and the effects being tested
  3. From still images to cinematic time
  4. Performance: voice, face, body and listening
  5. Continuity, memory and selective revision
  6. A method for showing possibilities without hiding failures
  7. The fragmented tool ecosystem
  8. Commercial aggregation and tool-conglomeration platforms
  9. Why an open and free alternative is necessary
  10. The Open Film Environment
  11. Evaluation, evidence and publication
    - 11.1 Global technology watch: capabilities, consolidation and volatility
  12. Labour, rights and adaptive governance
  13. Historical note: invention, independent routes and public support
  14. French and European perspectives
- Conclusion
- References

## 1. Purpose of the experiment

CineAI Lab asks a practical question: how far can an independent creator move from text, reference images and sound toward coherent dramatic cinema while retaining control over identity, performance, revision and export? The question is deliberately broader than image or video generation. It concerns the whole path by which an intention becomes a film.

The laboratory emerged under limited resources and with a preference for local or open tools. This constraint was productive. It made visible the difference between apparent access and real autonomy. A service can be easy to open yet difficult to integrate, expensive to repeat, impossible to audit or dependent on a model that may disappear from the platform. Conversely, an open repository can be legally available but practically unusable because of hardware requirements, version conflicts and installation complexity.

The aim of the first cycle was therefore not to claim that a feature film could already be produced by pressing one button. The aim was to build representative scenes, test specific effects, record failure honestly and identify the architecture needed for future work.

*The experiment is successful when it makes a capability or a limitation visible. It does not need to pretend that the entire film already exists.*

## **1.1 Human-AI collaboration: ChatGPT as a working partner**

CineAI Lab is also a record of sustained collaboration between a human creator and ChatGPT. The assistant was not used for one isolated prompt or one finished output. It became a conversational research and development partner across repeated cycles of writing, testing, failure analysis, redesign and publication.

The collaboration has covered screenplay and scene analysis; the formulation of cinematic test questions; visual ideation and image iterations; local Blender, Python and FFmpeg workflows; diagnosis of rendering and dependency failures; multilingual voice and viseme work; comparison of portrait-animation, video-generation and editing tools; website architecture and content; consolidation of the manuscript; and a continuing watch of technical, legal and institutional developments.

This accumulated process is itself one of the project's findings. A general AI assistant can connect domains that are usually separated between writing, software, image, sound, production research and documentation. It can preserve hypotheses, compare alternatives and turn failed attempts into the next test. At the same time, the work exposed a clear boundary: conversational intelligence does not by itself remove GPU constraints, incompatible repositories, proprietary services, missing APIs or the artistic weaknesses of current video engines.

The resulting authorship model is therefore explicit. ChatGPT contributes analysis, generation, synthesis and technical assistance. The protected narrative works, the decision to accept or reject outputs, the artistic direction and final responsibility remain human. CineAI Lab is independent and is not affiliated with OpenAI.

## **2. Creative demonstrators and the effects being tested**

Several original works were used as private test beds. Their plots are not reproduced here. What matters for the research is the set of cinematic effects and production demands they expose.

### **2.1 Restrained human performance**

One demonstrator centres on a formal public address delivered by an adult protagonist whose authority must coexist with fatigue, doubt and private tension. The experimental target is not the literal content of the speech. It is the preservation of a credible human presence across a sustained acted moment.

- Stable face, age, costume and environment across shot sizes.
- Natural multilingual voice with directed rhythm, pauses and breathing.
- Lip articulation integrated with jaw, cheeks, eyes, brows and head movement.
- Controlled gaze, hesitation, restraint and micro-expression.
- Listening and reaction shots that share the same emotional state.
- The ability to revise one line delivery without changing camera, identity or surrounding performances.

### **2.2 Non-human movement, scale and memory**

A second demonstrator tests animal locomotion, scale relationships, atmospheric continuity and transitions between present perception and remembered experience. The public research description stops at those effects; the protected narrative remains outside the document.

- Weight transfer, foot contact and coordination between trunk, head and body.
- Continuity of a herd or group under movement.

- Readable scale between very large and very small animals.
- Non-verbal expression of fear, attention, hostility, recognition and retreat.
- Transitions into memory that preserve who is observing, what is remembered and when the image occurs.
- Transformation or reflection effects that remain psychologically legible rather than becoming an unexplained change of creature.

## 2.3 Ensemble continuity and symbolic objects

A third demonstrator tests several adult characters whose relationships evolve through small shifts rather than spectacle. A recurring puzzle image functions as a continuity object and as the visual philosophy of the laboratory.

- Distinct adult identities maintained across domestic, social and professional settings.
- Slow changes in trust, withdrawal, suspicion, tenderness and disappointment.
- A recurring object whose physical state, missing part and visual meaning remain stable.
- Symbolic repetition without turning the motif into an explanatory cliché.
- Revision of one relationship beat without disturbing the wider ensemble arc.

Together, these demonstrators prevent the laboratory from optimizing for one easy case. A talking face does not test animal motion. A beautiful landscape does not test listening. A five-second clip does not test memory. The diversity of the works is therefore part of the method.

## 3. From still images to cinematic time

The earliest convincing outputs were still images. They were immediately useful for fictional casting, visual worlds, colour, composition, symbolic motifs and project communication. This success created a tempting assumption: if one image can already look cinematic, a sequence of similar images should become a film.

Time introduces a different order of difficulty. It adds anatomy, causality, rhythm, performance, spatial continuity and irreversible error. A hand that changes shape for one frame, a costume that drifts, a gaze that crosses the axis or a light that resets can destroy the credibility created by an otherwise beautiful image.

Blender and 2.5D compositing offered a controllable bridge. Layered images could be placed in depth, the camera could move and frames could be rendered locally. The pipeline succeeded technically, but the result frequently remained a pan, zoom or parallax effect. The frame moved while the character remained internally still.

This distinction became the central finding of the first cycle: movement is not action, and action is not performance. Camera motion can support a scene, but it cannot supply the intention of a character. The value of Blender therefore remains substantial for layout, compositing, camera logic, geometry, masks, inserts and export. It should not be mistaken for a complete actor engine.

## 4. Performance: voice, face, body and listening

Voice tests confirmed that intelligibility is only the first threshold. A dramatically useful voice needs identity, breath, pauses, accent, emotional direction and rights clarity. Smoothness can become a defect when a scene requires hesitation, fatigue or contradiction. Synthetic speech is already useful for timing, drafts, translation and some finished uses, but cinema requires a director to shape the performance rather than accept a generic delivery.

Viseme work transformed speech into structured timing data. A multilingual vocabulary could indicate mouth categories and connect phonetic events to an animation timeline. Yet the experiment also demonstrated that correct timing does not guarantee a believable face. If the lips are too small, the geometry mismatched or the cheeks and jaw remain frozen, the synchronized mouth becomes an attached graphic rather than part of a person.

Facial-animation systems can add blinks, head motion and lip sync, but their outputs must be judged as complete acted moments. The relevant question is not whether the mouth moved at the right syllable. It is whether the person appeared to think, listen, breathe and respond within the dramatic situation.

The performance stack is coupled:

- Vocal layer: tempo, volume, pitch contour, breath, hesitation, accent and intention.
- Facial layer: gaze, blink, tension, asymmetry, micro-expression and listening behaviour.
- Bodily layer: posture, gesture, weight, distance, orientation and stillness.
- Relational layer: attention, dominance, avoidance, attraction, fear and complicity.
- Temporal layer: preparation, action, reaction, recovery and silence.

A future system must preserve these layers as a performance score that can be revised. Natural-language direction remains important, but the intention should also survive as structured project data: what the character wants, what is hidden, when the gaze changes, where the silence begins and which choices are locked.

## 5. Continuity, memory and selective revision

The first experiments revealed that coherence breaks most often at tool boundaries. An image is resized for animation; an animation service changes the face; a voice file shifts the timing; a regenerated shot alters the costume; an editor receives a flattened clip with no record of the elements that produced it. Each handoff loses part of the film.

Long-form production therefore requires memory that is richer than a folder of previous frames. The system must distinguish at least four kinds of state:

- Canonical state: identity, biography, voice, world rules and approved design.
- Current state: clothing, location, injury, possession, weather and object condition.
- Relational state: trust, resentment, debt, attraction, fear and knowledge shared or withheld.
- Audience state: what has already been shown, promised or deliberately concealed.

Selective revision is equally important. Filmmaking advances through small changes: a softer line, a longer silence, a different weather condition, one missing prop, an alternate reaction. A system that can only regenerate the whole shot or scene turns every correction into a new risk. The desired environment must freeze approved dimensions and regenerate only the declared target.

*A generated shot should be an editable decision, not an irreversible event.*

## 6. A method for showing possibilities without hiding failures

The first cycle produced a simple experimental discipline. It is intended to reduce wasted work and separate marketing spectacle from observable capability.

1. Define one cinematic question: performance, locomotion, continuity, memory transition, object persistence or selective revision.
2. Prepare approved anchors for identity, setting, colour, sound and forbidden changes.
3. Build a short test, usually five to twelve seconds, rather than an undiagnosable long sequence.
4. Inspect the screen, not only the terminal log or exported file.
5. Record what was expected, what happened, the visible limitation and the next decision in ordinary prose.
6. Preserve representative failures as evidence instead of publishing only the best generation.
7. Assemble only shots that pass human review, using conventional editing when that remains the most reliable method.

The purpose is to demonstrate possibilities. A prototype may show that one aspect is ready while another remains fragile. This is more useful than presenting an apparently complete film whose success depends on

hidden retries, manual repairs and discarded outputs. CineAI Lab therefore treats unsuccessful results as part of the research record.

## 7. The fragmented tool ecosystem

No single tool currently owns the complete process. The experimental stack involved or considered several specialist families: image generation, compositing, camera control, speech synthesis, lip sync, portrait animation, locomotion, encoding, editing, sound and orchestration. Each family may contain open models, proprietary models, desktop applications, hosted notebooks and subscription services.

Open-source components offer inspectability, local control and the possibility of adaptation. They also create practical burdens: large downloads, incompatible versions, GPU assumptions, undocumented dependencies and continuous maintenance. A nominally free model may be inaccessible to a creator who lacks the required hardware or the ability to repair a Python environment.

Commercial services reduce installation friction, provide current models and often offer polished interfaces. Their limits are different: recurring credits, changing catalogues, opaque backends, account dependence, uncertain export lineage and terms that can change after a project has begun.

The problem is therefore not simply open versus commercial. It is the absence of a durable project layer that can survive the replacement of any one engine.

## 8. Commercial aggregation and tool-conglomeration platforms

A new category has become increasingly important: commercial platforms that conglomerate many creative tools or models behind one interface. The term is used here descriptively. These services may combine proprietary models, third-party models, editors, templates, enhancement, sound, agents and workflow management. Their value is real: they reduce setup time and allow users to compare capabilities without opening a separate account for every provider.

Higgsfield explicitly presents a multi-model video workspace, while Krea and Pollo AI similarly reduce the friction of moving among different generation systems. Adobe Firefly now also combines its own video model with partner models inside a familiar creative environment. API platforms such as fal.ai and Replicate expose large catalogues for developers. Magnific belongs to an adjacent category focused primarily on visual enhancement and image transformation; it illustrates specialised service consolidation but should not be treated here as equivalent to a broad multi-model video aggregator.

These platforms solve a genuine problem: tool discovery and access. They can provide one account, one credit system, one asset library and rapid availability of new models. For many users, this is more practical than maintaining several local repositories.

But aggregation is not the same as integration. A platform can place many models in the same menu without giving the film a persistent memory. It may simplify generation while leaving unresolved:

- stable characters and objects across multiple engines;
- the preservation of narrative and emotional state;
- selective regeneration with locked dimensions;
- editable project files rather than flattened outputs;
- transparent model versions and asset lineage;
- portable data when a model, price or platform changes;
- a production economy that does not require a growing stack of subscriptions.

The rapid success of aggregation services is therefore both an opportunity and a warning. It proves that creators want a coherent workspace. It also shows how easily the orchestration layer can become proprietary, subscription-based and difficult to leave.

## 9. Why an open and free alternative is necessary

CineAI Lab proposes that the public-interest response should not be another catalogue of links. It should be an integrated distribution: an open and free core environment that assembles specialist tools while preserving the project independently from them.

The word free requires precision. Advanced video generation consumes expensive computation, and no serious proposal can promise unlimited frontier-model inference at no cost. The achievable objective is different:

- The application, project format, orchestration logic and local workflows should be free and open source.
- Local models should be usable without subscription whenever hardware permits.
- Optional cloud compute should be replaceable and clearly priced rather than structurally required.
- A project should remain readable and editable if a provider disappears.
- Community-maintained presets should hide installation complexity without hiding provenance.
- Core capabilities should run on ordinary macOS, Windows and Linux systems, with graceful degradation on modest hardware.
- Rights, consent and model information should travel with the assets.

Such an environment would not attempt to defeat commercial platforms by copying every feature. It would provide the missing public infrastructure: durable project memory, open interchange, reproducible experiments and freedom to replace one model without rebuilding the film. Commercial engines could still be connected as optional modules, but they would no longer own the project.

The strategic contrast is therefore clear. Aggregators optimize access to tools. An open film environment must optimize continuity of intention.

## 10. The Open Film Environment

The consolidated proposal is an Open Film Environment built around the concepts developed in the earlier Open Film Model. It is not necessarily one neural network. It is a persistent, versioned project system in which different models and human departments cooperate without losing context.

Its principal modules would be:

- Project and story memory: scenes, causality, themes, unresolved questions and audience knowledge.
- Character memory: identity, voice, biography, relationships, knowledge, emotional trajectory and continuity locks.
- World and object memory: locations, physical state, recurring props, weather, institutions and visual language.
- Scene and shot specifications: purpose, blocking, performance, camera, sound, dependencies and permitted changes.
- Performance direction: subtext, gaze, breathing, gesture, listening, intensity and silence.
- Media synthesis connectors: interchangeable image, video, voice, motion and audio engines.
- Non-destructive editing: takes, timelines, branches, locked decisions and selective regeneration.
- Governance and provenance: sources, models, versions, permissions, human approvals and delivery status.

The primary unit is the project, not the prompt. Prompts become temporary instructions inside a larger object. This allows the system to answer production questions: Which version of the character is approved? Which scenes depend on this object? What does each person know here? What may change without affecting the rest of the cut?

An open interchange language would connect screenwriting, asset management, generation, editing and review. It should remain human-readable, permit vendor-specific extensions and retain stable identifiers when scenes are renamed or reordered.

## 11. Evaluation, evidence and publication

Evaluation should move from spectacle toward capability. Image quality and prompt adherence remain useful, but a film system must also be judged across time, revision and collaboration.

A concise capability profile can examine:

- identity persistence;
- world and object continuity;
- narrative memory;
- dramatic intention;
- performance direction;
- multilingual dialogue;
- camera and staging control;
- temporal coherence;
- selective regeneration and editability;
- provenance and human approval.

Results should declare the number of attempts, human intervention time, hardware or service tier, model versions and accepted defects. A representative failure should accompany the selected result. The goal is not to produce a universal beauty score, but to make the conditions of success visible and reproducible.

Human review remains mandatory. Directors, editors, performers and technical evaluators notice different failures. Their disagreement should be recorded rather than compressed into an opaque average. The final reference is whether the result serves the intended dramatic and aesthetic purpose.

### 11.1 Global technology watch: capabilities, consolidation and volatility

The technical landscape now changes fast enough that a fixed list of “best tools” would age before the manuscript. CineAI Lab therefore treats technology watch as part of the research method. The useful question is not which brand is temporarily ahead, but which capabilities are becoming dependable and which dependencies remain fragile.

Several developments in 2025–2026 illustrate the movement. Google introduced Veo 3.1 Lite as a lower-cost developer model. Kuaishou describes Kling 3.0 as a multimodal system combining text, image, audio and video with stronger consistency and shot-level control, while reporting expanding professional use in film, television, short drama, animation, advertising and gaming. Alibaba’s Wan 2.7 supports video continuation, audio synchronisation and multi-shot generation. Adobe Firefly increasingly combines generation, partner models and timeline editing in one environment. Runway’s own Gen-4.5 research simultaneously acknowledges persistent problems such as causal reasoning and object permanence. In China, the National Radio and Television Administration’s 2026 MediaAIAC programme explicitly calls for deployable AI applications across audiovisual content generation, story development, storyboarding, image and sound generation, digital humans, multilingual dubbing, provenance and safety. The significance is not that one region has solved cinema, but that model development, industrial deployment and public-sector coordination are moving on different tempos around the world.

The market is also moving through licensing and consolidation. The Disney–OpenAI agreement of December 2025 linked licensed characters, generative video, APIs and enterprise use of ChatGPT. Yet the subsequent discontinuation of the Sora web and app product in April 2026, with an API sunset announced for September, demonstrates the opposite lesson: even frontier interfaces can disappear quickly. A long-lived film project cannot make its memory, rights record or edit history dependent on one service.

These signals strengthen rather than weaken the open-film argument. Models will improve, prices will move and providers will change. The durable research object is the project layer that can compare engines, preserve approved decisions, document provenance and migrate when the surrounding market shifts.

## 12. Labour, rights and adaptive governance

The audiovisual sector is experiencing rational fear. Generative systems affect authorship, performance rights, employment, cost, content volume, distribution and public trust at the same time. Writers fear the conversion of their work into training material or unpaid drafts. Performers fear digital replicas and uses that exceed the original consent. Technicians face both automation and a growing volume of synthetic material that still requires correction and control. Producers see lower entry costs alongside uncertain rights and reputational risk.

Two responses would be equally damaging. One would allow deployment without consent, compensation, attribution or accountability. The other would attempt to freeze experimentation through legal uncertainty until the technological movement has relocated elsewhere. Rights should not become an excuse for invisibility, but regulation should not imitate an ostrich during an earthquake.

The practical middle requires governed experimentation:

- specific consent for voice, likeness, performance and protected assets;
- clear permitted uses, duration, territories and compensation;
- provenance from source material to final cut;
- worker participation before deployment decisions become irreversible;
- narrow pilots with the ability to stop, correct and revert;
- public reporting of benefits, defects and legal questions;
- support for independent creators, not only large studios and model providers.

European law already contains a useful mechanism: regulatory sandboxes intended to provide controlled environments for development, testing and regulatory learning. In the cultural sector, such spaces could connect legal advice, artistic experimentation, rights-holders and technical evaluation rather than forcing small projects to choose between paralysis and ungoverned risk.

## 13. Historical note: invention, independent routes and public support

The history invoked here is not a claim that cinema was invented by France alone. Moving-image technologies emerged through parallel experiments in several countries. France nevertheless played a decisive role at the moment when technical invention became collective projection, spectacle and then an artistic language. That history provides context for CineAI Lab, not a proof that artificial intelligence follows the same path.

### 13.1 From moving image to cinematic language

In 1895 the Lumière Cinématographe brought recording, printing and projection together in a compact apparatus. The Institut Lumière notes that the device allowed a film to become visible to an assembled audience, culminating in the commercial public screening of 28 December 1895 in Paris. The importance of this threshold was not simply mechanical: moving images became a shared public experience.

Georges Méliès then demonstrated how quickly a recording technology could be transformed into a technology of fiction. Drawing on stage illusion, optical tricks, camera stops, dissolves, superimpositions, colour and editing, he helped establish the imaginative grammar of cinematic effects. The Cinémathèque française describes him as a pioneer of cinematic spectacle and an inventor of film special effects.

The parallel with generative AI should remain modest. AI is not a new Cinématographe, and the industrial, legal and social conditions are radically different. The useful historical observation is narrower: a technical threshold acquires cultural meaning only when creators turn it into a language, when audiences can encounter the results, and when production methods become sufficiently accessible to support repeated experimentation.

## **13.2 Independent routes and artistic renewal**

French cultural history also contains repeated examples in which artistic renewal became visible partly by creating routes outside established selection systems. In 1863, the Salon des Refusés allowed the public to encounter works excluded from the official Salon. In 1874, Monet, Renoir, Morisot, Degas, Pissarro, Sisley, Cézanne and other artists chose to exhibit independently in the former studio of the photographer Nadar. The point is not that institutional recognition is unnecessary. It is that new artistic languages may need a public space before existing categories are ready to receive them.

These episodes should not be romanticised as a simple conflict between artists and institutions. Their relevance lies in a recurring cultural mechanism: new forms may need spaces in which they can be seen and discussed before established categories know how to receive them. Independent circulation can create visibility; institutions can later provide continuity, preservation, professional recognition and resources.

## **13.3 In and Off: coexistence rather than opposition**

Avignon offers another French model. Jean Vilar founded the Festival d'Avignon in 1947 to reach a young and new public with theatre different from the Parisian habits of the time. In 1967, Vilar still described Avignon as a place of research and creation. The independent Off, born around the Festival in 1966, later developed a parallel ecology of presentation. The relevance here is not to copy a festival structure, but to note how institutional selection and independent emergence can coexist around the same cultural field.

For CineAI Lab, this is a useful historical image rather than a ready-made organisational proposal. Innovation may need both selective institutional spaces and independent zones of experimentation. The strength lies in circulation between them, not in treating one as virtuous and the other as an obstacle.

Post-war French cinema also developed a strong tradition of public support. The CNC was created in 1946 to support the creation and economy of cinematic works, and the advance on receipts, launched in 1960, was explicitly designed to renew creation, encourage first films and support independent cinema whose artistic ambition could not necessarily find equilibrium under market conditions alone. French cultural history therefore combines invention, independent initiative and public support rather than offering a single model of innovation.

# **14. French and European perspectives**

The present transition should be understood against that longer history. Generative AI does not automatically democratise filmmaking, but it lowers some barriers between an idea and its first audiovisual expression. It can make visual development, voice experiments, previsualisation, short movement studies and rapid variation available to creators who do not command the resources of a studio. At the same time, dependence on closed platforms, compute costs, fragmented subscriptions and uncertain rights can reproduce new forms of concentration.

## **14.1 France: a historical and contemporary junction**

France enters this transition with unusual assets: a globally recognised cinema culture, a dense production and festival ecosystem, public mechanisms for supporting creation, strong traditions of authors' rights, specialised schools and research institutions, and a national technology sector that now includes major AI ambitions. France 2030 and Ministry of Culture programmes already support digital transformation, immersive culture and the responsible appropriation of artificial intelligence by cultural and creative industries. This does not yet

amount to a unified strategy for AI-native cinema, but it shows that the technical, cultural and public-policy components are beginning to converge.

The CNC's 2026 programme supporting exploratory uses of generative AI is another sign that experimentation is entering the institutional landscape. Four structures were selected to host artists and authors in exploratory residencies. These initiatives do not yet amount to a unified audiovisual AI strategy, but they show that artistic experimentation, technical change and public support are beginning to meet.

The Mistral case remains relevant in this context, but not as a grievance. Mistral is a private company whose development has also been associated with public investment, national technological-sovereignty ambitions and institutional partnerships. CineAI Lab's attempts to establish a substantive exchange had not succeeded by early August 2026. This experience is retained only as an illustration of a broader coordination question: significant technological capacity can exist in the same national ecosystem as independent cultural experimentation without an obvious bridge between them.

The French perspective is therefore not primarily to claim historical ownership of a new technology. It is to recognise that France already possesses several of the cultural, legal, artistic and technical ingredients required to examine how AI can widen access to filmmaking without dissolving authorship, performance, professional competence or public trust.

## **14.2 Europe: diversity, rights and interoperability**

The European perspective is wider. Europe is not a single audiovisual market or a single language space. Its cinema has long depended on multilingual production, cross-border circulation and coproduction. Eurimages, operating since 1989, illustrates this cooperative tradition by supporting independent film across a large group of European states and partner countries. This diversity can be a constraint for technology deployment, but it is also an important test of whether generative systems can serve different languages, legal traditions and cultural expectations.

European regulation is now becoming operational at the same time as creative adoption accelerates. The transparency obligations of Article 50 of the AI Act began to apply on 2 August 2026, including requirements concerning AI-generated or manipulated content. The European Audiovisual Observatory has emphasised that AI in the audiovisual sector intersects with copyright, personality rights, data protection, media regulation and other legal frameworks. Regulation therefore forms part of the production environment rather than a separate debate that can be postponed until after the tools mature.

At the same time, the European Commission's Apply AI Strategy explicitly includes the cultural, creative and media sectors and frames AI adoption around the movement from laboratory to market, the connection of infrastructure, data and testing facilities, and the development of trustworthy European capacity. Europe is therefore attempting to combine adoption with safeguards rather than choosing between technological acceleration and legal restraint.

The strategic question is whether this combination can become practically usable for creators and small production structures. Europe is unlikely to reproduce the concentration of capital and platforms found in the United States or China. It may instead have an opportunity to develop a different strength: interoperable and multilingual production methods in which rights, provenance, cultural diversity and human editorial control are designed into the workflow rather than added after generation.

## **14.3 An open European horizon**

CineAI Lab should be situated within this horizon rather than as a finished institutional proposal. Its experiments suggest that the central problem is no longer simply access to a powerful video model. It is the continuity of the creative project across writing, image, movement, voice, performance, editing, rights and revision.

France provides a credible place from which to pursue this inquiry because of its cinematic history and cultural infrastructure. Europe provides the larger scale, linguistic plurality, legal framework and tradition of

collaboration needed to test whether an open film environment can remain portable across borders and providers.

*The historical opportunity is not to automate cinema, but to widen the number of people who can enter its workshop while preserving the memory, rights and human intention that make the result a film.*

## Conclusion

The first CineAI Lab draft contained several documents because the project was discovering its own subject. The consolidated view is now simpler: the research is about preserving a cinematic intention across tools, time and revision. The path to that conclusion matters. It was built through sustained work with ChatGPT across narrative analysis, image development, experimental design, Blender/Python/FFmpeg troubleshooting, voice and viseme tests, comparison of open and commercial tools, website architecture, documentation, source verification and global technology watch. The value of that collaboration is cumulative rather than spectacular: ideas were repeatedly confronted with outputs, technical constraints, failed attempts and revised hypotheses. The protected stories, acceptance criteria and final creative decisions remain human.

Current systems already offer remarkable possibilities in writing assistance, visual development, voice, short movement, enhancement and rapid variation. The global watch shows that video models are becoming more multimodal, better integrated and cheaper to access, while licensing agreements and creative-suite aggregation are expanding. It also shows volatility: services, interfaces and commercial terms can change or disappear faster than a film project can be developed.

The missing piece is an open production environment in which the project remains larger than any model. It must remember characters and objects, carry dramatic direction, permit selective change, record provenance and keep human approval visible. The demonstrators are intended to show that this direction is possible and sufficiently concrete to justify a coordinated next phase involving technical, cultural, legal and financial partners.

*Technology can generate an image. Cinema begins when human intention remains visible from one shot, one tool and one revision to the next.*

## References

All links below were checked against official sources on 7 August 2026. The list is selective: it supports the claims made in this manuscript and the technology-watch examples rather than attempting a complete market survey.

### Human-AI collaboration and OpenAI

OpenAI, “What is ChatGPT: FAQ”. [Official source ↗](#)

OpenAI, “What to know about the Sora discontinuation”. [Official source ↗](#)

OpenAI, “The Walt Disney Company and OpenAI reach landmark agreement...”, 11 December 2025. [Official source ↗](#)

### Global audiovisual AI watch

Google, “Build with Veo 3.1 Lite”, 31 March 2026. [Official source ↗](#)

Adobe Firefly, “Generate videos using text prompts”, updated June 2026. [Official source ↗](#)

Runway, “Introducing Runway Gen-4.5” and documented limitations. [Official source ↗](#)

Kuaishou Technology, “Kling AI Launches 3.0 Model”, 5 February 2026. [Official source ↗](#)

Alibaba Cloud Model Studio, “Wan Image-to-Video 2.7”, updated 15 July 2026. [Official source ↗](#)

National Radio and Television Administration of China, Fifth MediaAIAC AI Application Innovation Competition, 21 July 2026. [Official source ↗](#)

Kuaishou Technology, Fourth Quarter and Full Year 2025 Results: professional Kling AI applications in film, television and animation. [Official source ↗](#)

## Aggregation and tool-conglomeration services

Higgsfield, AI Video Generator: multi-model video workspace. [Official service page ↗](#)

Krea, AI Video Generator: multiple video models in one subscription. [Official service page ↗](#)

Magnific, official platform page. Cited as an adjacent image-enhancement service rather than a broad video-model aggregator. [Official service page ↗](#)

Pollo AI, AI Video: access to multiple leading video models in one interface. [Official source ↗](#)

fal, generative media platform: large catalogue of image, video, audio and 3D models through APIs. [Official source ↗](#)

Replicate, official models: maintained generative models exposed through stable APIs. [Official source ↗](#)

## France and Europe

French Ministry of Culture, strategy for cultural and responsible AI, 2 July 2025. [Official source ↗](#)

CNC, four selected exploratory generative-AI residencies, 21 April 2026. [Official source ↗](#)

French Ministry of Culture, France 2030 call for digital transition and appropriation of AI by cultural and creative industries, 16 May 2025. [Official source ↗](#)

Mistral AI, company mission and approach. [Official source ↗](#)

Bpifrance, Mistral AI Series C of €1.7 billion, 9 September 2025. [Official source ↗](#)

Bpifrance, MGX, Mistral AI and NVIDIA, French AI campus joint venture, 19 May 2025. [Official source ↗](#)

Bpifrance and Mistral AI, strategic deployment of Le Chat, 18 May 2026. [Official source ↗](#)

French Ministry of Culture / CSPLA, report on the copyright status of AI productions, 16 July 2026. [Official source ↗](#)

European Commission, Article 50 transparency guidelines, published 20 July 2026; obligations apply from 2 August 2026. [Official source ↗](#)

European Commission, Apply AI webinar for cultural, creative and media sectors, June 2026. [Official source ↗](#)

European Audiovisual Observatory, “AI in the Audiovisual Sector: Navigating the Current Legal Landscape”. [Council of Europe source ↗](#)

## Historical perspective

Institut Lumière, “Le Cinématographe Lumière”. [Official institutional source ↗](#)

La Cinémathèque française, Georges Méliès exhibition and museum materials. [Official institutional source ↗](#)

Musée d’Orsay, “Paris 1874: Inventing Impressionism”. [Official institutional source ↗](#)

Festival d’Avignon, “Histoire et archives”. [Official source ↗](#)

Festival Off Avignon, “60 ans d’histoire”. [Official source ↗](#)

Festival d'Avignon, archive 1967: Avignon described as a place of research and creation. [Official source ↗](#)

CNC, “Les dates-clés du CNC”. [Official source ↗](#)

CNC, “Avance sur recettes avant réalisation”. [Official source ↗](#)

## Document status

Version 0.9 · August 2026. Consolidated working manuscript prepared from the documented CineAI Lab experimental record. This revision makes the sustained collaboration with ChatGPT explicit, extends the global audiovisual-AI watch, adds the Chinese institutional and industrial perspective, corrects the treatment of aggregation services, and provides verified clickable official-source links for the claims and examples retained in the manuscript.